

# HW 3

## Integer Sort

Philip Nelson

2018 September 21

### Introduction

The purpose of this assignment is to write an MPI program that performs a parallel integer sort. It is written in the "master/slave" style. The process with id zero is the master; it generates an array of random integers. The master then sends chunks of the array to the slave processes to be sorted. When a slave receives a chunk, it sorts the data and returns it to the master. Once the master has sent out all the data, it waits to receive the sorted chunks from the slaves. As the sorted chunks are received, they are merged with the sorted array. Once all the chunks have been received and merged, the program reports the time that it took to sort the list and exits.

### Code

```
1  #include "random.hpp"
2  #include <algorithm>
3  #include <chrono>
4  #include <iostream>
5  #include <mpi.h>
6  #include <numeric>
7  #include <vector>
8
9  #define MCW MPI_COMM_WORLD
10
11 enum Tag
12 {
13     UNSORTED,
14     SORTED
15 };
16
17 std::vector<int> merge(std::vector<int> const& a, std::vector<int>
    const& b)
18 {
19     std::vector<int> merged;
```

```

20  unsigned int i = 0, j = 0;
21  while (i < a.size() && j < b.size())
22  {
23      if (a[i] < b[j])
24      {
25          merged.push_back(a[i++]);
26      }
27      else
28      {
29          merged.push_back(b[j++]);
30      }
31  }
32
33  std::copy(begin(a) + i, end(a), std::back_inserter(merged));
34  std::copy(begin(b) + j, end(b), std::back_inserter(merged));
35
36  return merged;
37  }
38
39  void random_fill(std::vector<int>::iterator b,
40                  std::vector<int>::iterator e,
41                  int low = 0,
42                  int high = 1000)
43  {
44      std::for_each(b, e, [&](int& a) { a = randInt(low, high); });
45  }
46
47  int main(int argc, char** argv)
48  {
49      int rank, world_size;
50
51      MPI_Init(NULL, NULL);
52      MPI_Comm_rank(MCW, &rank);
53      MPI_Comm_size(MCW, &world_size);
54
55      if (0 == rank)
56      {
57          int n = 11;
58          if (argc >= 2)
59          {
60              n = std::stoi(argv[1]);
61          }
62
63          /* ----- */
64          /* Generate Unsorted Data */
65          /* ----- */

```

```

66     std::vector<int> unsorted;
67     unsorted.resize(n);
68     random_fill(std::begin(unsorted), std::end(unsorted));
69
70     auto start = std::chrono::high_resolution_clock::now();
71     /* ----- */
72     /* Send Unsorted Data Chunks */
73     /* ----- */
74     const int chunksize = n / (world_size - 1);
75     for (auto i = 0; i < world_size - 2; ++i)
76     {
77         MPI_Send((begin(unsorted) + (chunksize * i)).base(),
78                 chunksize,
79                 MPI_INT,
80                 i + 1,
81                 Tag::UNSORTED,
82                 MCW);
83     }
84
85     MPI_Send((begin(unsorted) + (chunksize * (world_size - 2))).base(),
86             (chunksize + (n % chunksize)),
87             MPI_INT,
88             world_size - 1,
89             Tag::UNSORTED,
90             MCW);
91
92     /* ----- */
93     /* Receive Sorted Data */
94     /* ----- */
95     MPI_Status stat;
96     std::vector<int> data;
97     std::vector<int> result;
98     int size;
99     for (auto i = 0; i < world_size - 1; ++i)
100    {
101        MPI_Probe(MPI_ANY_SOURCE, Tag::SORTED, MCW, &stat);
102        MPI_Get_count(&stat, MPI_INT, &size);
103        data.resize(size);
104        MPI_Recv(data.data(),
105                size,
106                MPI_INT,
107                MPI_ANY_SOURCE,
108                Tag::SORTED,
109                MCW,
110                MPI_STATUS_IGNORE);
111

```

```

112     sorted = merge(sorted, data);
113 }
114 auto end = std::chrono::high_resolution_clock::now();
115 auto total_time =
116     std::chrono::duration<double, std::milli>(end - start).count();
117     std::cout << "Time: " << total_time << " ms\n";
118 }
119 else
120 {
121     int n;
122     MPI_Status stat;
123     MPI_Probe(0, 0, MCW, &stat);
124     MPI_Get_count(&stat, MPI_INT, &n);
125
126     std::vector<int> data;
127     data.resize(n);
128
129     MPI_Recv(data.data(), n, MPI_INT, 0, Tag::UNSORTED, MCW,
130             MPI_STATUS_IGNORE);
131     std::sort(begin(data), end(data));
132     MPI_Send(data.data(), n, MPI_INT, 0, Tag::SORTED, MCW);
133 }
134 MPI_Finalize();
135
136 return EXIT_SUCCESS;
137 }

```

## Output

```
# mpic++ main.cpp -O3 -o release.out
```

```
# mpiexec -n 4 release 100000
```

```
Time: 9.75266 ms
```

## Findings

I ran this program on lists of random integers from 100 to 1,000,000,000 elements and compared its performance to `std::sort` running on a single thread. The results are detailed in Figure 1. The graph displays the ratio of MPI Sort to `std::sort`. Lists upto 10,000 elements are sorted faster by `std::sort` but between 10,000 and 100,000,000 elements, the MIP Sort was faster. After 100,000,000 elements however, `std::sort` takes over as the faster sort.

My theory for this behavior is that sending the messages between threads becomes too expensive between 100,000,000 and 1,000,000,000 elements and `std::sort` is able to run faster with a single thread.

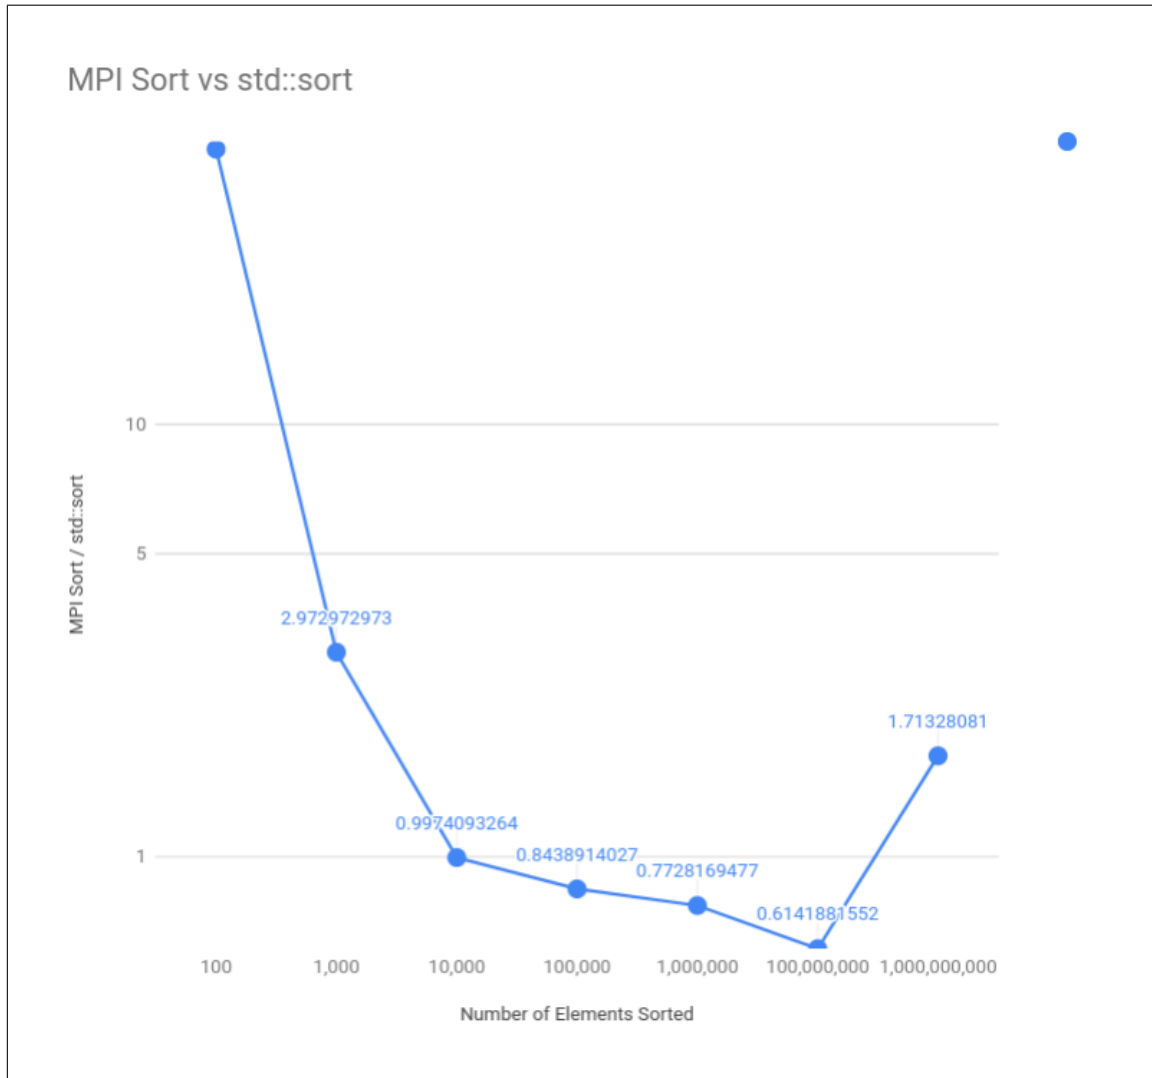


Figure 1: MPI Sort vs std::sort