

HW 7

Global Sum

Philip Nelson

2018 October 12

Introduction

The purpose of this assignment is to write an MPI program that performs a global sum. I have implemented several different variants. The first takes advantage of the cube network communication. Second uses ring communication. Third is a naive master slave method. Fourth uses the building MPI_Allgather function. I made each iteration sleep for one tenth of a second in order to simulate some kind of work being executed.

Code

```
1  #include "random.hpp"
2  #include <algorithm>
3  #include <iomanip>
4  #include <iostream>
5  #include <mpi.h>
6  #include <unistd.h>
7  #include <vector>
8
9  #define slep 100000
10
11 void printlper(int data, std::string title = "")
12 {
13     int rank;
14     int word_size;
15
16     MPI_Comm_rank(MPI_COMM_WORLD, &rank);
17     MPI_Comm_size(MPI_COMM_WORLD, &word_size);
18
19     if (0 == rank)
20     {
21         int* dArray = new int[word_size];
```

```

22     MPI_Gather(&data, 1, MPI_INT, dArray, 1, MPI_INT, 0, MPI_COMM_WORLD
23         );
24     std::cout << title << '\n';
25     for (int i = 0; i < word_size; ++i)
26     {
27         std::cout << std::setw(5) << i << std::setw(5) << dArray[i] << "\
28             n";
29     }
30     std::cout << std::endl;
31 }
32 else
33 {
34     MPI_Gather(&data, 1, MPI_INT, nullptr, 1, MPI_INT, 0,
35         MPI_COMM_WORLD);
36 }
37 int cube(int c, int sendData, int rank)
38 {
39     int recvData;
40     auto dest = rank ^ (1 << c);
41
42     MPI_Send(&sendData, 1, MPI_INT, dest, 0, MPI_COMM_WORLD);
43     MPI_Recv(&recvData, 1, MPI_INT, dest, 0, MPI_COMM_WORLD,
44         MPI_STATUS_IGNORE);
45
46     return recvData;
47 }
48 int ring(int dir, int sendData, int rank, int world_size)
49 {
50     int recvData;
51     auto dest = (rank + 1 * dir) % world_size;
52     auto src = (rank - 1 * dir) % world_size;
53
54     MPI_Send(&sendData, 1, MPI_INT, dest, 0, MPI_COMM_WORLD);
55     MPI_Recv(&recvData, 1, MPI_INT, src, 0, MPI_COMM_WORLD,
56         MPI_STATUS_IGNORE);
57
58     return recvData;
59 }
60 void cubeSum(int num, int rank, int world_size)
61 {
62     int log2n = log2(world_size);

```

```

63     for (auto i = 0; i < log2n; ++i)
64     {
65         num += cube(i, num, rank);
66         usleep(slep);
67     }
68     printf("cube sum\n");
69 }
70
71 void ringSum(int num, int rank, int world_size)
72 {
73     int next, prev = num;
74
75     for (auto i = 0; i < world_size - 1; ++i)
76     {
77         next = ring(1, prev, rank, world_size);
78         num += next;
79         prev = next;
80         usleep(slep);
81     }
82     printf("ring sum\n");
83 }
84
85 void masterSlaveSum(int num, int rank, int world_size)
86 {
87     if (0 == rank)
88     {
89         int recvData;
90         for (auto i = 1; i < world_size; ++i)
91         {
92             MPI_Recv(&recvData,
93                     1,
94                     MPI_INT,
95                     MPI_ANY_SOURCE,
96                     0,
97                     MPI_COMM_WORLD,
98                     MPI_STATUS_IGNORE);
99             num += recvData;
100             usleep(slep);
101         }
102         printf("master slave sum\n    0    " << num << "\n\n");
103     }
104     else
105     {
106         MPI_Send(&num, 1, MPI_INT, 0, 0, MPI_COMM_WORLD);
107     }
108 }

```

```

109
110 void mpiAllReduce(int num)
111 {
112     MPI_Allreduce(&num, &num, 1, MPI_FLOAT, MPI_SUM, MPI_COMM_WORLD);
113     usleep(slep);
114     printlper(num, "all reduce");
115 }
116
117 int main(int argc, char** argv)
118 {
119     MPI_Init(&argc, &argv);
120
121     int rank, world_size;
122     MPI_Comm_rank(MPI_COMM_WORLD, &rank);
123     MPI_Comm_size(MPI_COMM_WORLD, &world_size);
124
125     if (0 != (world_size & (world_size - 1)))
126     {
127         if (rank == 0)
128         {
129             std::cerr << "There must be a power of 2 number of threads\n";
130         }
131
132         MPI_Finalize();
133         exit(EXIT_SUCCESS);
134     }
135
136     int num;
137     if (0 == rank)
138     {
139         std::vector<int> data(world_size);
140         random_double_fill(begin(data), end(data), 0, 10);
141         MPI_Scatter(data.data(), 1, MPI_INT, &num, 1, MPI_INT, 0,
142                     MPI_COMM_WORLD);
143         printlper(num, "original data");
144     }
145     else
146     {
147         MPI_Scatter(nullptr, 1, MPI_INT, &num, 1, MPI_INT, 0,
148                     MPI_COMM_WORLD);
149         printlper(num);
150     }
151
152     auto t1 = MPI_Wtime();
153     cubeSum(num, rank, world_size);
154     auto t2 = MPI_Wtime();

```

```

153  ringSum(num, rank, world_size);
154  auto t3 = MPI_Wtime();
155  masterSlaveSum(num, rank, world_size);
156  auto t4 = MPI_Wtime();
157  mpiAllReduce(num);
158  auto t5 = MPI_Wtime();
159
160  if (0 == rank)
161      std::cout << "cube: " << t2 - t1 << "\nring: " << t3 - t2
162          << "\nmaster slave: " << t4 - t3 << "\nall reduce: " <<
            t5 - t4
163          << "\n";
164
165  MPI_Finalize();
166
167  return (EXIT_SUCCESS);
168  }

```

Output

```
# mpic++ -O3 main.cpp -o release.out
```

```
# mpiexec -n 8 --oversubscribe release.out
```

```
original data
```

0	5
1	3
2	0
3	8
4	4
5	8
6	3
7	9

```
cube sum
```

0	40
1	40
2	40
3	40
4	40
5	40
6	40
7	40

```
ring sum
```

0	40
1	40
2	40
3	40
4	40
5	40
6	40
7	40

```
master slave sum
```

0	40
---	----

```
all reduce
```

0	40
1	40
2	40
3	40
4	40

5	40
6	40
7	40

```
cube: 0.305294
ring: 0.708875
master slave: 0.701991
all reduce: 0.101125
```

Findings

The cube sum was the best performing of my own sorting functions which is unsurprising since it only needs to compute $\log_2 n$ sums where n is the world size. Therefore it is about 43% faster. The ring sum was as performant as the master slave sum however only the master process is left with the sum. If you want all processes to have the global sum, the ring sum would be a better option. Getting a good time estimate for the reduce / all reduce isn't possible using this setup but I imagine it is the best option for performing global sums.