

HW 8

Load Balancing

Philip Nelson

2018 November 2

Introduction

The purpose of this assignment is to implement sender-initiated distributed random dynamic load balancing with white/black ring termination. I started each process with one task, the tasks were doubles between $[1, 3)$ and processing a task mean sleeping for as many milliseconds as the value of the task times 1000. Each task was assignment a maximum number of tasks to produce between $[1024, 2048]$. If a process ever had more than 16 tasks to process, it would send the next two tasks to a random process. If it ever sent tasks to a process less than it's rank, it would turn it's token black. When process 0 finished it's tasks, it would send a white token to the next process. When other processes finished, they would receive the token and send it on to the next process turning it black if their own token was black. They would then wait to receive an action from process zero. Process zero would tell each process to continue or finalize depending on what color token it received. When a white token returned to process zero, all processes would finalize and report their work.

Code

The code is broken up into two files, main.cpp and random.hpp. The files are included below.

0.1 main.cpp

```
1 #include "random.hpp"
2 #include <chrono>
3 #include <iostream>
4 #include <mpi.h>
5 #include <queue>
6 #include <thread>
7
8 enum TAG
9 {
10     NEW_WORK,
```

```

11     TOKEN,
12     ACTION,
13     CONTINUE,
14     FINALIZE
15 };
16
17 enum TOKEN
18 {
19     WHITE,
20     BLACK
21 };
22
23 void process(double task)
24 {
25     int work = (int)(task*1000);
26     std::this_thread::sleep_for(std::chrono::milliseconds(work));
27 }
28
29 void getNewWork(std::queue<double>& tasks, int& workRecv)
30 {
31     int flag;
32     MPI_Status stat;
33     do
34     {
35         MPI_Iprobe(MPI_ANY_SOURCE, TAG::NEW_WORK, MPI_COMM_WORLD, &flag, &
36             stat);
37         if (flag)
38         {
39             double task;
40             MPI_Recv(&task,
41                 1,
42                 MPI_DOUBLE,
43                 stat.MPI_SOURCE,
44                 stat.MPI_TAG,
45                 MPI_COMM_WORLD,
46                 MPI_STATUS_IGNORE);
47             tasks.push(task);
48             ++workRecv;
49         }
50     } while (flag);
51
52 void makeMoreWork(std::queue<double>& tasks,
53     int& workMade,
54     int maxWorkMade,
55     int workMin,

```

```

56             int workMax)
57 {
58     if (workMade < maxWorkMade)
59     {
60         auto newWork = random_int(1, 3);
61         for (auto i = 0; i < newWork; ++i)
62         {
63             tasks.push(random_double(workMin, workMax));
64             ++workMade;
65         }
66     }
67 }
68
69 int main(int argc, char** argv)
70 {
71     MPI_Init(&argc, &argv);
72
73     int rank, world_size;
74     MPI_Comm_rank(MPI_COMM_WORLD, &rank);
75     MPI_Comm_size(MPI_COMM_WORLD, &world_size);
76
77     const int workMin = 1, workMax = 3, maxWorkLoad = 16, amntToSend = 2,
78             maxWorkMade = random_int(1024, 2048);
79     int workMade = 0, workDone = 0, workSent = 0, workRecv = 0;
80     enum TOKEN token = TOKEN::WHITE;
81
82     std::queue<double> tasks;
83     if (rank == 0)
84     {
85         tasks.push(random_double(workMin, workMax));
86
87         do // while token == black
88         {
89             token = TOKEN::WHITE;
90             while (!tasks.empty())
91             {
92                 process(tasks.front());
93                 tasks.pop();
94                 ++workDone;
95
96                 /* Create more work */
97                 makeMoreWork(tasks, workMade, maxWorkMade, workMin, workMax);
98
99                 /* Check for incoming work */
100                getNewWork(tasks, workRecv);
101

```

```

102      /* Send work to other processes */
103      if (tasks.size() > maxWorkLoad)
104      {
105          for (auto i = 0; i < amntToSend; ++i)
106          {
107              int dest = random_int(rank + 1, rank + world_size - 1) %
                  world_size;
108              MPI_Send(&tasks.front(),
109                      1,
110                      MPI_DOUBLE,
111                      dest,
112                      TAG::NEW_WORK,
113                      MPI_COMM_WORLD);
114              tasks.pop();
115              ++workSent;
116          }
117      }
118  }
119
120      /* Send Token */
121      MPI_Send(
122          &token, 1, MPI_INT, rank + 1 % world_size, TAG::TOKEN,
123          MPI_COMM_WORLD);
124
125      std::cout << rank << " -- send token -> "
126          << (token == TOKEN::BLACK ? "BLACK" : "WHITE") << '\n';
127
128      /* Receive Token */
129      MPI_Recv(&token,
130              1,
131              MPI_INT,
132              world_size - 1,
133              TAG::TOKEN,
134              MPI_COMM_WORLD,
135              MPI_STATUS_IGNORE);
136
137      std::cout << rank << " <- recv token -- "
138          << (token == TOKEN::BLACK ? "BLACK" : "WHITE") << '\n';
139
140      if (token == TOKEN::BLACK)
141      {
142          enum TAG action = TAG::CONTINUE;
143          for (auto i = 1; i < world_size; ++i)
144          {
145              MPI_Send(&action, 1, MPI_INT, i, TAG::ACTION, MPI_COMM_WORLD)
146              ;
147          }
148      }

```

```

145     }
146 }
147
148 } while (token == TOKEN::BLACK);
149
150 /* Kill all processes */
151 enum TAG action = TAG::FINALIZE;
152 for (auto i = 1; i < world_size; ++i)
153 {
154     MPI_Send(&action, 1, MPI_INT, i, TAG::ACTION, MPI_COMM_WORLD);
155 }
156 }
157 else // SLAVE
158 {
159     tasks.push(random_double(workMin, workMax));
160     enum TAG action;
161     do // while action == continue
162     {
163         token = TOKEN::WHITE;
164         while (!tasks.empty())
165         {
166             process(tasks.front());
167             tasks.pop();
168             ++workDone;
169
170             /* Create more work */
171             makeMoreWork(tasks, workMade, maxWorkMade, workMin, workMax);
172
173             /* Check for incoming work */
174             getNewWork(tasks, workRecv);
175
176             /* Send work to other processes */
177             if (tasks.size() > maxWorkLoad)
178             {
179                 for (auto i = 0; i < amntToSend; ++i)
180                 {
181                     int dest = random_int(rank + 1, rank + world_size - 1) %
182                         world_size;
183                     if (dest < rank) token = TOKEN::BLACK;
184                     MPI_Send(&tasks.front(),
185                             1,
186                             MPI_DOUBLE,
187                             dest,
188                             TAG::NEW_WORK,
189                             MPI_COMM_WORLD);
190                     tasks.pop();

```

```

190         ++workSent;
191     }
192 }
193 }
194
195 enum TOKEN recvToken;
196 /* Receive Token */
197 MPI_Recv(&recvToken,
198         1,
199         MPI_INT,
200         rank - 1,
201         TAG::TOKEN,
202         MPI_COMM_WORLD,
203         MPI_STATUS_IGNORE);
204
205 if (recvToken == TOKEN::BLACK)
206 {
207     token = TOKEN::BLACK;
208 }
209 std::cout << rank << " <- recv token -- "
210         << (recvToken == TOKEN::BLACK ? "BLACK" : "WHITE") << '
211         \n';
212
213 /* Send Token */
214 MPI_Send(&token,
215         1,
216         MPI_INT,
217         (rank + 1) % world_size,
218         TAG::TOKEN,
219         MPI_COMM_WORLD);
220
221 std::cout << rank << " -- send token -> "
222         << (token == TOKEN::BLACK ? "BLACK" : "WHITE") << '\n';
223
224 MPI_Recv(
225     &action, 1, MPI_INT, 0, TAG::ACTION, MPI_COMM_WORLD,
226     MPI_STATUS_IGNORE);
227
228 } while (action == TAG::CONTINUE);
229
230 std::cout << rank << " -- FINALIZE -- " << '\n';
231 std::cout << rank << " -- work done -- " << workDone << '\n';
232
233 int totalWorkDone;

```

```

233     MPI_Allreduce(&workDone, &totalWorkDone, 1, MPI_INT, MPI_SUM,
234                 MPI_COMM_WORLD);
235     double percentWorkDone = (double)workDone / totalWorkDone * 100;
236     std::cout << rank << "-- percent done -- " << percentWorkDone << '\n'
237         ;
238     if (rank == 0) std::cout << "Total Work Done " << totalWorkDone << '\n';
239     MPI_Finalize();
240
241     return EXIT_SUCCESS;
242 }

```

0.2 random.hpp

```

1  #ifndef RANDOM_HPP
2  #define RANDOM_HPP
3
4  #include <algorithm>
5  #include <functional>
6  #include <random>
7
8  /**
9   * Generate a random number from [low, high]
10  *
11  * @param low The lower bound
12  * @param high The upper bound
13  * @return A random number on the range [low, high]
14  */
15 int random_int(int low, int high)
16 {
17     static std::random_device rd;
18     static std::mt19937 mt(rd());
19     std::uniform_int_distribution<> dist(low, high);
20     return dist(mt);
21 }
22
23 /**
24  * Generate a random number from [low, high]
25  *
26  * @param low The lower bound
27  * @param high The upper bound
28  * @return A random number on the range [low, high]
29  */
30 double random_double(double low, double high)
31 {

```

```

32     static std::random_device rd;
33     static std::mt19937 mt(rd());
34     std::uniform_real_distribution<> dist(low, high);
35     return dist(mt);
36 }
37
38 /**
39  * Fill a container from [first, last) with random numbers from [low,
40  *   high]
41  * @param first Iterator to beginning of range to fill
42  * @param last  Iterator to end of range to fill
43  * @param low   The lower bound
44  * @param high  The upper bound
45  */
46 template <typename it>
47 void random_int_fill(it first, it last, const int low, const int high)
48 {
49     static std::random_device rd;
50     static std::mt19937 mt(rd());
51     std::uniform_int_distribution<> dist(low, high);
52     std::generate(first, last, std::bind(dist, mt));
53 }
54
55 /**
56  * Fill a container from [first, last) with random numbers from [low,
57  *   high)
58  * @param first Iterator to beginning of range to fill
59  * @param last  Iterator to end of range to fill
60  * @param low   The lower bound
61  * @param high  The upper bound
62  */
63 template <typename it>
64 void random_double_fill(it first, it last, const double low, const
65     double high)
66 {
67     static std::random_device rd;
68     static std::mt19937 mt(rd());
69     std::uniform_real_distribution<double> dist(low, high);
70     std::generate(first, last, std::bind(dist, mt));
71 }
72 #endif

```


Output

```
# mpic++ -O3 main.cpp -o release.out
```

```
# mpiexec -n 4 release.out
```

```
0 -- send token -> WHITE
1 <- recv token -- WHITE
1 -- send token -> BLACK
2 <- recv token -- BLACK
2 -- send token -> BLACK
3 <- recv token -- BLACK
3 -- send token -> BLACK
0 <- recv token -- BLACK
0 -- send token -> WHITE
1 <- recv token -- WHITE
1 -- send token -> WHITE
2 <- recv token -- WHITE
2 -- send token -> WHITE
3 <- recv token -- WHITE
3 -- send token -> WHITE
0 <- recv token -- WHITE
```

```
0 -- FINALIZE --
0 -- work done -- 900
```

```
1 -- FINALIZE --
1 -- work done -- 839
```

```
2 -- FINALIZE --
2 -- work done -- 974
```

```
3 -- FINALIZE --
3 -- work done -- 1202
```

```
Total Work Done 3915
0-- percent done -- 22.9885
1-- percent done -- 21.4304
2-- percent done -- 24.8787
3-- percent done -- 30.7024
```

Findings

From the report at the end of the program execution, each process processed between 20% and 30% of the total tasks. This means that despite producing a different amount of tasks, they were able to divide the tasks between themselves fairly well.