

Homework 3 - Count Sort

Philip Nelson

12th October 2021

1

When parallelizing the outer loop, the variables `i`, `j`, and `count` should be private and `a`, `n`, and `temp` should be shared.

2

The call to `memcpy` can not be parallelized as it is written. This would result in the sorted array being copied in its entirety by each thread. Instead, each thread can copy a chunk of the data in the following manner.

```
const int rank = omp_get_thread_num();
const int n_per_thread = n / thread_count;
const int start = rank * n_per_thread;
const int my_n = (rank == thread_count - 1) ? n - (rank * n_per_thread) : n_per_thread;
memcpy(a + start, temp + start, my_n * sizeof(int));
```

It is important to only free `temp` in one thread by putting it outside of the parallel context.

3

I implemented two versions of the parallel count sort algorithm. The first is a modern c++ implementation based on iterators, the second parallelizes the code given in the assignment description.

This project includes a CMakeLists.txt with three projects: `count_sort`, `count_sort_benchmark` and `tests`. The project `count_sort` requires command line arguments for number of threads, data count, and optionally a seed for the pseudo random number generator (default 100). It generates a sequence of integers, prints the original sequence, and then the sorted sequence to stdout. The pseudo random number generator, `std::rand`, was seeded with 100 for each trial. The following commands will build and run the `count_sort` project.

```
mkdir build
cd build
cmake -DCMAKE_BUILD_TYPE=Release ..
make count_sort
./count_sort/count_sort 4 25
```

```
# [ 16, 2, 5, 10, 7, 24, 6, 21, 7, 22, 6, 4, 23, 23, 16, 2, 6, 14, 11, 19, 20, 17, 13, 16, 7 ]
# [ 2, 2, 4, 5, 6, 6, 6, 7, 7, 7, 10, 11, 13, 14, 16, 16, 16, 17, 19, 20, 21, 22, 23, 23, 24 ]
```

Timing results

Unsurprisingly this algorithm has a time complexity of $O(n^2)$ which can clearly be seen in Figure 1.

Question 2 asked to parallelize the copying of the sorted temporary array. I found that this portion of the algorithm was taking approximately 0.001% of the total runtime. After parallelization, performance suffered slightly, see Figure

2 (my version) and Figure 3 (given code). The difference is so little it likely falls within the margin of error. These tests were run with a 20 trial average on 10,000 elements.

I also experimented with the static scheduling chunksize. Figure 4 and Figure 5 show timings results for 8 threads on 10,000 elements for chunk sizes from 1 to 100 and 1 to 1,000 by 10 respectively. It's interesting to see how some chunksizes are significantly worse and how the trend does not increase monotonically.




Figure 1: Exponential Runtime Complexity




Figure 2: Parallelization of outer loop and move

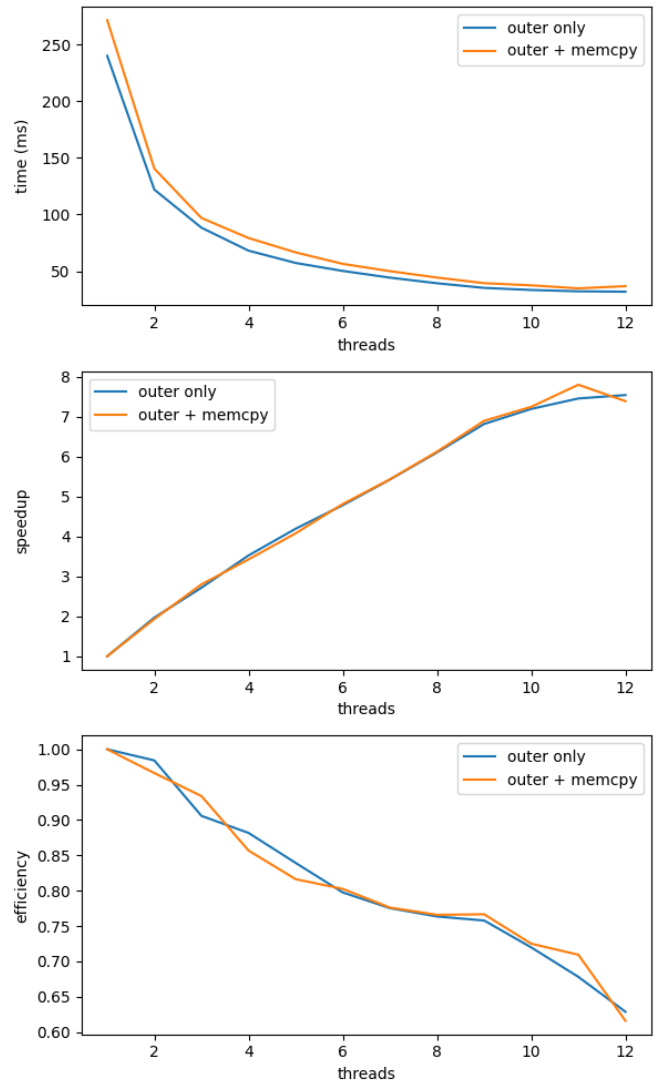


Figure 3: Parallelization of outer loop and memcpy




Figure 4: Static Scheduling with Chunksize 1-100




Figure 5: Static Scheduling with Chunksize 1-1000