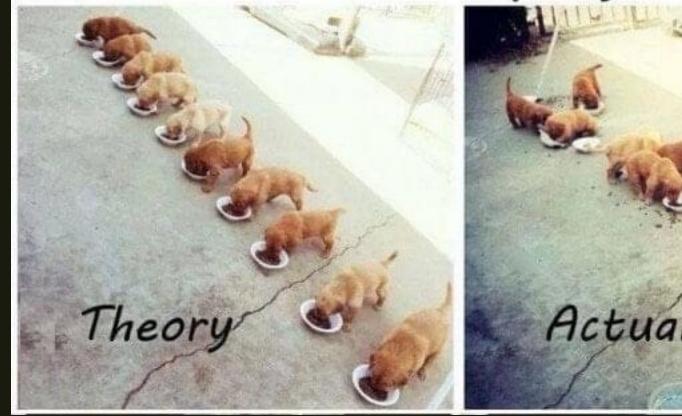
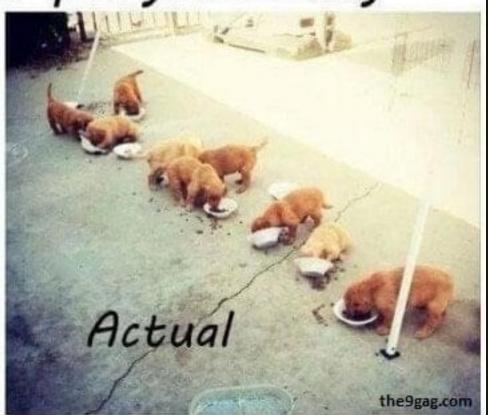
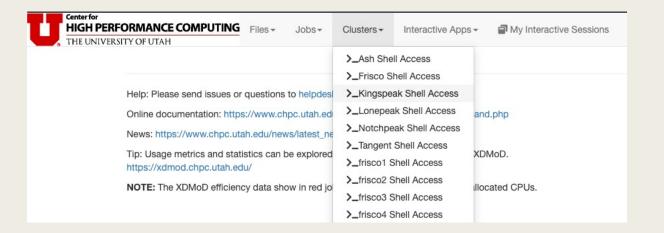
## Multithreaded programming





## Preliminary test for CHPC access

- Try to login to ondemand CHPC: <a href="https://ondemand.chpc.utah.edu/">https://ondemand.chpc.utah.edu/</a>
- Try to start a terminal on one of the clusters ("Shell Access")



## Setting up a Makefile to build your program

```
#target1 [target2 ...]: [pre-req-1 pre-req-2 ...]
#
     [command1
#
      command2
#
      .....]
all: hello
hello: hello.o
         gcc hello.o -o hello -lpthread
hello.o: pth_hello.c
        gcc -c pth_hello.c -o hello.o
```

## Run Make

- Linux/MacOS
  - make
- Windows
  - mingw32-make



# CRITICAL SECTIONS

## Estimating π

$$\pi = 4\left(1 - \frac{1}{3} + \frac{1}{5} - \frac{1}{7} + \dots + (-1)^n \frac{1}{2n+1} + \dots\right)$$

```
double factor = 1.0;
double sum = 0.0;
for (i = 0; i < n; i++, factor = -factor) {
    sum += factor/(2*i+1);
}
pi = 4.0*sum;</pre>
```

How can we parallelize this code using multiple threads?

#### A thread function for computing $\pi$

```
void* Thread_sum(void* rank) {
  long my_rank = (long) rank;
  double factor;
  long long i;
  long long my n = n/thread count;
   long long my first i = my n*my rank;
  long long my last i = my first i + my n;
   if (my_first_i \% 2 == 0) /* my_first_i is even */
     factor = 1.0;
   else /* my_first_i is odd */
      factor = -1.0;
  for (i = my_first_i; i < my_last_i; i++, factor = -factor) {</pre>
      sum += factor/(2*i+1);
   return NULL;
   /* Thread_sum */
```

## Using a dual core processor

	n						
	$10^{5}$	$10^{6}$	107	$10^{8}$			
π	3.14159	3.141593	3.1415927	3.14159265			
1 Thread	3.14158	3.141592	3.1415926	3.14159264			
2 Threads	3.14158	3.141480	3.1413692	3.14164686			

Note that as we increase n, the estimate with one thread gets better and better.

### Possible race condition

```
y = Compute(my_rank);

x = x + y;
```



Time	Thread 0	Thread 1	
1	Started by main thread		
2	Call Compute ()	Started by main thread	
3	Assign $y = 1$	Call Compute()	
4	Put x=0 and y=1 into registers	Assign $y = 2$	
5	Add 0 and 1	Put x=0 and y=2 into registers	
6	Store 1 in memory location x	Add 0 and 2	
7		Store 2 in memory location x	

## **Busy-Waiting**

- A thread repeatedly tests a condition, but, effectively, does no useful work until the condition has the appropriate value.
- Beware of optimizing compilers, though!

```
y = Compute(my_rank);
while (flag != my_rank);
x = x + y;
flag++;
```

flag initialized to 0 by main thread

#### Pthreads global sum with busy-waiting

```
void* Thread_sum(void* rank) {
   long my rank = (long) rank;
   double factor;
   long long i;
   long long my_n = n/thread_count;
   long long my_first_i = my_n*my_rank;
   long long my_last_i = my_first_i + my_n;
   if (my first i \% 2 == 0)
      factor = 1.0;
   else
      factor = -1.0;
   for (i = my_first_i; i < my_last_i; i++, factor = -factor) {</pre>
      while (flag != my_rank);
      sum += factor/(2*i+1);
      flag = (flag+1) \% thread count;
   return NULL:
  /* Thread_sum */
```

Busy wait every loop iteration!

Can we do a little better?

#### Global sum function with critical section after loop

```
void* Thread_sum(void* rank) {
   long my_rank = (long) rank;
   double factor, my_sum = 0.0;
   long long i;
   long long my_n = n/thread_count;
   long long my_first_i = my_n*my_rank;
   long long my_last_i = my_first_i + my_n;
   if (my_first_i \% 2 == 0)
      factor = 1.0;
   else
      factor = -1.0;
  for (i = my_first_i; i < my_last_i; i++, factor = -factor)</pre>
     my_sum += factor/(2*i+1);
  while (flag != my_rank);
  sum += my_sum;
  flag = (flag+1) % thread_count;
  return NULL;
  /* Thread_sum */
```

- A thread that is busy-waiting may continually use the CPU accomplishing nothing.
- Mutex (mutual exclusion) is a special type of variable that can be used to restrict access to a critical section to a single thread at a time.



- Used to guarantee that one thread "excludes" all other threads while it executes the critical section.
- The Pthreads standard includes a special type for mutexes: pthread\_mutex\_t.

■ When a Pthreads program finishes using a mutex, it should call

```
int pthread_mutex_destroy(pthread_mutex_t* mutex_p /* in/out */);
```

■ In order to gain access to a critical section a thread calls

```
int pthread_mutex_lock(pthread_mutex_t* mutex_p /* in/out */);
```

■ When a thread is finished executing the code in a critical section, it should call

```
int pthread_mutex_unlock(pthread_mutex_t* mutex_p /* in/out */);
```

#### C++ mutex

std::mutex my\_mutex;

```
my_mutex.lock()
//lock the mutex
```

my\_mutex.unlock() //unlock the mutex

How can we use mutex in our "PI" program?

#### Global sum function that uses a mutex (1)

```
void* Thread_sum(void* rank) {
   long my_rank = (long) rank;
   double factor;
   long long i;
   long long my_n = n/thread_count;
   long long my_first_i = my_n*my_rank;
   long long my_last_i = my_first_i + my_n;
   double my sum = 0.0;
   if (my first i \% 2 == 0)
     factor = 1.0;
   else
     factor = -1.0;
  for (i = my first i; i < my last i; i++, factor = -factor) {
     my sum += factor/(2*i+1);
  pthread mutex lock(&mutex);
  sum += my_sum;
  pthread_mutex_unlock(&mutex);
  return NULL:
  /* Thread_sum */
```

Threads	Busy-Wait	Mutex
1	2.90	2.90
2	1.45	1.45
4	0.73	0.73
8	0.38	0.38
16	0.50	0.38
32	0.80	0.40
64	3.56	0.38

$$\frac{T_{\rm serial}}{T_{\rm parallel}} \approx {\tt thread\_count}$$

Run-times (in seconds) of  $\pi$  programs using n = 108 terms on a system with two four-core processors.

		Thread					
Time	flag	0	1	2	3	4	
0	0	crit sect	busy wait	susp	susp	susp	
1	1	terminate	crit sect	susp	busy wait	susp	
2	2		terminate	susp	busy wait	busy wait	
:	:			:	:	:	
?	2			crit sect	susp	busy wait	

Possible sequence of events with busy-waiting and more threads than cores.



#### Issues

- Busy-waiting enforces the order threads access a critical section.
- Using mutexes, the order is left to chance and the system.
- There are applications where we need to control the order threads access the critical section.

#### Problems with a mutex solution

Example: product of matrices is not commutative

```
/* n and product_matrix are shared and initialized by the main thread */
/* product_matrix is initialized to be the identity matrix */
void* Thread_work(void* rank) {
   long my_rank = (long) rank;
   matrix_t my_mat = Allocate_matrix(n);
   Generate_matrix(my_mat);
   pthread_mutex_lock(&mutex);
   Multiply_matrix(product_mat, my_mat);
   pthread_mutex_unlock(&mutex);
   Free_matrix(&my_mat);
   return NULL;
} /* Thread_work */
```

#### Another example: a first attempt at sending messages using pthreads

```
/* messages has type char**. It's allocated in main. */
/* Each entry is set to NULL in main.
void *Send_msq(void* rank) {
  long my_rank = (long) rank;
  long dest = (my_rank + 1) % thread_count;
  long source = (my_rank + thread_count - 1) % thread_count;
  char* my_msg = malloc(MSG_MAX*sizeof(char));
   sprintf(my_msg, "Hello to %ld from %ld", dest, my_rank);
  messages[dest] = my_msg;
   if (messages[my_rank] != NULL)
      printf("Thread %ld > %s\n", my_rank, messages[my_rank]);
   else
      printf("Thread %ld > No message from %ld\n", my_rank, source);
  return NULL:
  /* Send_msg */
```

What can we do to ensure that every thread receives its message?

#### Use mutex??

```
1 ...
2 pthread_mutex_lock(mutex[dest]);
3 ...
4 messages[dest] = my_msg;
5 pthread_mutex_unlock(mutex[dest]);
6 ...
7 pthread_mutex_lock(mutex[my_rank]);
8 printf("Thread %ld > %s\n", my_rank, messages[my_rank]);
9 ...
```

#### What can happen here?

#### Syntax of the various semaphore functions

```
#include <semaphore.h>
```

Semaphores are not part of Pthreads; you need to add this.

```
int sem_init(
    sem_t* semaphore_p /* out */,
    int shared /* in */,
    unsigned initial_val /* in */);
```

```
int sem_destroy(sem_t* semaphore_p /* in/out */);
int sem_post(sem_t* semaphore_p /* in/out */);
int sem_wait(sem_t* semaphore_p /* in/out */);
```

Post: increments the semaphore value Wait, will block if the semaphore is 0, otherwise it will decrement the semaphore value and continue... and... decrement when unblock!

## Use of semaphore

```
/* messages is allocated and initialized to NULL in main */
/* semaphores is allocated and initialized to 0 (locked) in
     main */
void* Send_msg(void* rank) {
  long my_rank = (long) rank;
  long dest = (my_rank + 1) % thread_count;
  char* my_msg = malloc(MSG_MAX*sizeof(char));
  sprintf(my_msq. "Hello to %ld from %ld". dest. my_rank):
  messages[dest] = my_msg;
  sem_post(&semaphores[dest])
        /* "'Unlock" the semaphore of dest */
  /* Wait for our semaphore to be unlocked */
  sem_wait(&semaphores[my_rank]);
  printf("Thread %1d > %s\n", my_rank, messages[my_rank]);
  return NULL:
  /* Send_msg */
```



#### Barriers

- Synchronizing the threads to make sure that they all are at the same point in a program is called a barrier.
- No thread can cross the barrier until all the threads have reached it.

#### Using barriers to time the slowest thread

```
/* Shared */
double elapsed_time;
/* Private */
double my_start, my_finish, my_elapsed;
Synchronize threads;
Store current time in my_start;
/* Execute timed code */
Store current time in my_finish;
my_elapsed = my_finish - my_start;
elapsed = Maximum of my_elapsed values;
```

## Using barriers for debugging

```
point in program we want to reach;
barrier;
if (my_rank == 0) {
   printf("All threads reached this point\n");
   fflush(stdout);
}
```

## Busy-waiting and a Mutex

- Implementing a barrier using busy-waiting and a mutex is straightforward.
- We use a shared counter protected by the mutex.
- When the counter indicates that every thread has entered the critical section, threads can leave the critical section.

## Busy-waiting and a Mutex

```
/* Shared and initialized by the main thread */
int counter; /* Initialize to 0 */
int thread count;
                                         We need one counter
pthread_mutex_t barrier_mutex;
                                         variable for each
                                         instance of the barrier,
void* Thread_work(. . .) {
                                         otherwise problems
                                         are likely to occur.
   /* Barrier */
   pthread_mutex_lock(&barrier_mutex);
   counter++:
   pthread_mutex_unlock(&barrier_mutex);
   while (counter < thread_count);</pre>
   . . .
```

#### Implementing a barrier with semaphores

```
/* Shared variables */
int counter; /* Initialize to 0 */
sem_t count_sem; /* Initialize to 1 */
sem t barrier sem; /* Initialize to 0 */
void* Thread_work(...) {
   /* Barrier */
   sem wait(&count sem);
   if (counter == thread_count-1) {
     counter = 0:
      sem post(&count sem);
      for (j = 0; j < thread_count -1; j++)
        sem_post(&barrier_sem);
     else {
     counter++:
      sem_post(&count_sem);
      sem wait(&barrier sem);
```